


Minimum-Sufficient Payloads Reduce Context Without Lowering Judged Answer Quality



Quality Signals Preserved

LLM-as-judge: 9.2/10 (excellent)

Payload verifier: 13/13 views passed

Fail-soft fallbacks: 0

Real Q1 end-to-end run. Token counts compare the legacy full-dossier-style prompt payload against the compiled minimal-sufficient PayloadView